

10-718 Assignment 1: Problem Proposal

Yixin Dong and Akaash R. Parthasarathy

Problem Description

LLMs are increasingly capable of writing complete system code

- LLMs run on serving engines   etc.
- Serving engines call GPU kernels
- GPU kernels are highly specialized
 - For both hardware and the actual workload
 - Require significant effort and expertise to optimize

Questions:

- Can we generate kernels that are correct and achieve strong performance on different workloads from the trace distribution?
- What kind of LLM/agent design is suitable for this task?
- Model efficiency (reducing benchmark calls and kernel generation latency)

FlashInfer Trace Schema

- Use FlashInfer Trace Schema to study data-driven GPU kernel generation
 - Kernel **Definition**
 - **Workload** specification
 - Generate candidate **Solutions**
 - Use FlashInfer-Bench to run correctness and performance benchmarking and record results as new **Traces**

Target Kernels

Fused MoE

Fused Mixture-of-Experts kernels with FP8 support.

FP8 MoE

Sparse Attention

Deepseek Sparse Attention from [Deepseek V3.2](#)

Indexer

Attention


Gated Delta Net

Gated Delta Net used in [Qwen3-Next](#)

GDN Decode

GDN Prefill

Collection Process



Collection Methods




We have three curation paths (increasing levels of automation):

- Manually inject dump statements
 - in the execution path
- Auto-tracing via `flashinfer_bench.tracing`
 - needs adapter → schema
- LLM-assisted generation
 - Claude skills to write dumps/adapters

Overview of the Collected Workloads

Workloads:

- dsa_paged: 151
- gdn: 120
- moe: 19



Proposed Agent Design

Prompt: describe mathematical definition, basic guidelines

N Iterative Feedback Rounds -> best performing kernel

Generate Triton and CUDA kernels for B200

Baseline: Let LLMs directly generate a CUDA kernel with minimal prompt design

Input :Definition, Language, Hardware

Output : Sol^* (best solution)

$\mathcal{S} \leftarrow \emptyset$

$Agent \leftarrow \text{CodeAgent.Initialize(Definition, Language, Hardware)}$

$Sol_0 \leftarrow \text{Agent.Generate()}$

for $i \leftarrow 0$ **to** $N - 1$ **do**

$Trace_i \leftarrow \text{FlashInfer-Bench.Benchmark(Definition, Sol_i)}$

if $Trace_i.Status = \text{PASSED}$ **then**

$\mathcal{S} \leftarrow \mathcal{S} \cup \{(Sol_i, Trace_i)\}$

$Sol_{i+1} \leftarrow \text{Agent.Optimize}(Trace_i)$

$Sol^* \leftarrow \arg \max_{(Sol_i, Trace_i) \in \mathcal{S}} Trace_i.Speedup$

return Sol^*

Thanks!